




NeoMundi

The Fundamental AI Behavioral Measurement Layer

A common metrology infrastructure for multiple uses
and heterogeneous AI systems

One layer · Multiple uses · Heterogeneous infrastructures

Measure AI behavior.
Feed every use case.
Connect every infrastructure.

NeoMundi builds the fundamental layer for measuring AI behavior. One measurement can feed multiple uses — governance, observability, audit, compliance — and articulate with heterogeneous infrastructures, without replacing your existing tools.

AI behavior is variable. It must become measurable.

Organizations today are building extensively on AI systems whose real behavior — stability, coherence, drift — is not yet covered by a common measurement layer. NeoMundi builds that fundamental layer, connecting the AI's real behavior to operational decision-making.

Measure

Observe the model's or agent's behavior during execution.

Govern

Turn observations into signals the organization's policy can act on.

Prove

Link observation, rule, decision and action in a reconstructible trace.

**NeoMundi measures. Your policy decides.
Every decision stays traceable and verifiable.**

One measurement layer. Multiple uses and infrastructures.

What NeoMundi provides

- objective runtime measurements
- actionable behavioral signals
- auditable governance artefacts
- interfaces for third-party systems
- longitudinal knowledge
- reconstructible proofs

Can articulate with it

- GRC platforms
- compliance platforms
- observability tools
- auditors
- certifiers
- insurers
- risk teams
- AI integrators
- agentic systems
- orchestration platforms

These infrastructures keep their function and independence: they consume the measurement, they are not replaced by it.

AI no longer just answers. It acts.

The risk shifts from the final output to the authorization to act.

AI systems can now:

- call tools
- trigger workflows
- modify data
- coordinate other agents
- send communications
- initiate actions with limited human oversight

« Is the answer correct? »

becomes

« Should the system be allowed to continue or act in its current state? »

Response

Decision

Tool call

Action

Operational
effect

A governance policy cannot work without runtime signals

Organizations already have:

- models
- human validations
- compliance frameworks
- internal rules
- monitoring tools
- technical logs

But an independent layer able to observe the system's behavioral state while it is running is often missing.

Without this layer, the organization knows an execution took place. It does not always know:

when to intervene

why an action was authorized

which signal triggered an escalation

how to reconstruct the decision months later

A governance policy without runtime measurement remains an intention.

AI can be sovereign in where it is hosted and still be uncontrolled in its behavior

Hosting AI in Europe, on European infrastructure and under European law, is an essential condition of sovereignty.

But that is not enough.

Without continuous runtime measurement, an organization controls where its AI runs, but not necessarily:

its drifts

its instabilities

its semantic variations

its coherence issues

its factual risks

their operational consequences

***Sovereignty without runtime metrology
is blind sovereignty.***

The measurement layer that makes AI behavior observable and exploitable

NeoMundi develops the missing metrology layer between the model's behavior and the operational decision. Our infrastructure makes it possible to:

- continuously measure stability
- measure coherence
- detect semantic variations
- detect drifts and breaks
- track factual risk
- observe latency and consumption
- identify regime changes
- produce traceable, auditable runtime proofs
- feed governance actions
- connect signals to compliance and execution systems

True sovereignty means controlling where AI runs, how it behaves, and what it is authorized to trigger.

Measuring the real behavior of models and agents

Stability

Observe the regularity and stability of behavior during execution.

Coherence

Identify losses of semantic or structural coherence.

Semantic variations

Detect significant changes in meaning, wording or trajectory.

Drifts

Identify evolutions, breaks or regime changes over time.

Factual risk

Flag outputs requiring stronger factual validation.

Latency and tokens

Measure runtime behavior while accounting for cost and performance.

Regimes

Identify stable, transitional or unstable states.

Trace

Link each measurement to its execution context.

A signal is not a verdict.

It is an operational input for governance.

- 1 Observation during execution
- 2 Production of measurements
- 3 Signal detection
- 4 Application of the client's policy
- 5 Governed action
- 6 Creation of a verifiable trace

Possible actions

continue	verify	regenerate
reroute	escalate	validate
limit	stop	trace

NeoMundi provides the measurement and the signals. The organization keeps control of the decision.

Observe what can still be corrected. Govern what cannot.

OBS

Governance observability

For responses or outputs that can still be corrected.

- baseline
- longitudinal tracking
- comparison
- alerts
- history
- trace
- drift analysis
- semantic analysis

Full privacy

NeoMundi can receive only the necessary metrics, signals, snapshots or artefacts. Prompts and responses can stay within the client's perimeter depending on the architecture chosen.

GOV

Governance during execution

For actions that are hard to correct after execution. ControlTower measures within the execution path and produces a signal before the workflow's next step.

- continue
- regenerate
- reroute
- validate
- limit
- stop
- escalate

No retention

Prompts can pass through in real time with no durable storage logic. The retention configuration depends on the integration mode and the client's policy.

Central criterion: can the action still be corrected after execution?

One single layer instead of several fragmented controls

Before NeoMundi

- technical monitoring
- factuality
- drift detection
- auditability
- human oversight
- vendor comparison
- proof
- orchestration
- compliance

These functions are often spread across several tools, teams and processes.

With NeoMundi

A common layer brings together:

- runtime observations
- behavioral signals
- semantic variations
- alerts
- comparison
- history
- proofs
- action interfaces

NeoMundi does not replace your governance: it gives it runtime signals, proof, and the ability to act at the right moment.

The pilot measures value instead of promising it abstractly

Human oversight

Human time per case, escalation rate, share of genuinely useful reviews, reduction in undifferentiated checks.

Detection

Time between signal and intervention, earlier drift detection, reduced reaction time.

Runtime costs

Tokens generated, avoidable generations, regenerations, latency, cost per acceptable output.

Audit and proof

Reconstitution time, artefact quality, reduced manual work, proof availability.

Consolidation of building blocks

Number of blocks replaced, consolidated interfaces, reduced manual coordination.

Compliance readiness

Risk management, traceability, human control, continuous monitoring, incident investigation.

Pilot output: baseline, observed signals, intervention rules, and measured impact on oversight, cost, risk and traceability.

Who does NeoMundi create the most value for?

Financial and fiduciary decisions

Banking, insurance, credit, fraud, compliance, transaction validation.

Sensitive professional workflows

Healthcare, legal, public sector, human resources, critical services.

Autonomous and multi-agent systems

Agents connected to tools, APIs, databases and third-party systems.

AI integrators and platforms

A premium governance, oversight and proof layer they can embed in their own offerings.

Risk, compliance and audit teams

Access to actionable technical signals and artefacts for their processes.

Start where an error has a measurable consequence.

Preparing compliance during operation

ControlTower supports the technical implementation of several key obligations by producing measurements and proofs during operation. NeoMundi does not replace legal analysis, regulatory classification, certification, or the competent authority.

Article 9 — Risk management system

Measured signals, thresholds, alerts and longitudinal history.

Article 12 — Record-keeping

Identifiers, timestamps, logs and traceability artefacts.

Article 14 — Human oversight

Actionable signals to guide a review or a human intervention.

Article 26 — Deployer obligations

Monitoring of operation in the field.

Article 72 — Post-market monitoring

Continuous collection and analysis of relevant data across the lifecycle.

Article 73 — Reporting of serious incidents

Technical and historical elements that can be used to qualify and investigate an incident.

GDPR: data minimization · storage limitation · accountability · data protection by design

An architecture designed for multiple governance frameworks

NeoMundi is preparing a progressive mapping of its capabilities against the main regulatory, standards-based and internal frameworks.

- EU AI Act

- ISO/IEC 42001

- NIST AI RMF

- GDPR

- Internal company frameworks

- Sector-specific requirements

- Other international regulations to be integrated progressively

Do not claim that NeoMundi automatically certifies compliance with these frameworks.

NeoMundi produces measurements, signals and technical proofs that can support the implementation, documentation and audit of governance policies.

An instrument backed by continuous observation

The NeoMundi Observatory documents AI behaviors over time and feeds ControlTower's maturity. It publishes:

- weekly barometers
- protocols
- longitudinal analyses
- de-identified public data
- behavioral cartographies
- named studies
- methodological reports

Its purpose is not to produce a commercial ranking. It is to make visible behaviors that would otherwise remain unmeasured.

Each observation:

- improves the instrument
- enriches the knowledge base
- makes it possible to compare regimes
- feeds drift analyses
- improves the reading of semantic issues
- strengthens the ability to govern AI over time

FLAG signals correlated with confirmed problematic outputs in around 76% of cases across a cumulative corpus.

Result observed under the conditions studied, not a universal guarantee.

Technology ecosystem

infomaniak

The Ethical Cloud 

The NeoMundi Observatory is developed with the support of Infomaniak, a Swiss infrastructure partner committed to digital sovereignty, privacy and a responsible cloud.



NeoMundi is a member of the NVIDIA Inception program, dedicated to startups developing advanced artificial intelligence technologies.


The NeoMundi client pays for a complete, measured observation

1 NeoMundi observation = one fully instrumented execution

An observation includes:

- | | |
|----------------|---------------------------|
| 1 the request | 2 the associated response |
| 3 the protocol | 4 the runtime capture |
| 5 the analysis | 6 the comparison |
| 7 the signals | 8 the history |
| 9 the proof | 10 the ability to act |

The client does not pay for a simple API request. They pay for a complete, instrumented observation that can be used for:



Pricing tiers, wallets, offers and advanced analytics will be presented in the commercial annex.

Adoption program

The first 100 clients and partners

Facilitate initial adoption and quickly document NeoMundi's operational value.

The program can include:

- preferential access
- guided onboarding
- selection of a pilot workflow
- initial allocation of observations
- creation of a baseline
- definition of useful signals
- design of an intervention policy
- close follow-up
- results report
- preferential terms for early participants
- possibility of joint publication or case study, with the client's agreement

The goal is not just to connect an API. It is to demonstrate measurable operational value.

One workflow. One baseline. One measurable policy.

1 Select a workflow with real stakes.

2 Establish the behavioral baseline.

3 Identify useful signals.

4 Define the intervention rules.

5 Measure operational value.

Expected outputs


- baseline
- intervention rules
- impact on cost
- impact on traceability
- observed signals
- impact on oversight
- impact on risk
- deployment recommendations

One layer. Multiple uses.

The adoption of agentic AI requires a new operational layer:

- measure the system's state during execution
- understand how its behavior evolves
- define what it is authorized to do
- intervene at the right moment
- keep proof of the controls applied

Measure AI behavior.
Feed every use case.
Connect every infrastructure.



NeoMundi

The fundamental AI behavioral measurement layer.

contact@neomundi.io · controltower.neomundi.io · neomundi.io