



NeoMundi

The missing measurement context for AI behavior at runtime

NeoMundi measures the behavioral state of an AI system at a specific point in time, within a defined measurement framework.

ONE API CALL · NO INFRASTRUCTURE REPLACEMENT · PRIVACY-FIRST · BYOK

Your system.

Your decisions.

Our measurement signal.

This independent, timestamped and comparable signal provides the context needed to interpret observations, detections, audits, diagnoses, benchmarks, insurance assessments and evidence records — without replacing the infrastructures that produce them.

THE PROBLEM

Observing an AI is no longer enough to understand its behavior.

Detecting, auditing, diagnosing, comparing, insuring or documenting an AI system requires knowing what state it was in at the moment in question. This context data does not yet exist in an independent, timestamped and comparable form.

| A technical log shows that an execution took place — not what behavioral state the system was in.

| A one-off human validation says nothing about the system's regularity over time.

| A compliance rule defines a framework — it does not measure whether the system adhered to it at a given moment.

| An incident reveals a problem after the fact — rarely the signal that would have allowed it to be anticipated.

Without this context, observation, audit or evidence rest on partial elements.

The real behavior of models and agents

NeoMundi observes the behavior of a model or agent during execution, within a defined measurement framework. Each axis is measured within explicit limits — not as an absolute evaluation of response quality.

Stability

Observe the regularity of behavior during execution.

Semantic variations

Detect significant changes in meaning, wording or trajectory.

Factual risk

Flag outputs requiring stronger factual validation.

Regimes

Identify stable, transitional or unstable states.

Coherence

Identify losses of semantic or structural coherence.

Drifts

Identify evolutions, breaks or regime changes over time.

Latency and tokens

Measure runtime behavior while accounting for cost and performance.

Trace

Link each measurement to its execution context.

A stability signal is not a judgment of correctness. The two are measured separately.

A structured, timestamped, comparable and traceable signal

Each measurement produces a signal a third-party system can act on — not a closed conclusion.

Objective runtime measurements

Produced during execution, within the defined protocol.

Auditable governance artefacts

Designed to be reviewed and verified after the fact.

Longitudinal knowledge

Comparison of the same system over time.

Actionable behavioral signals

Transmittable to a third-party policy or infrastructure.

Interfaces for third-party systems

Designed for consumption by existing tools.


Reconstructible proofs

Each measurement links observation, context and timestamp.

A signal is not a verdict. It is context data that your infrastructure interprets within its own framework.

One integration. Multiple lightweight applications.

One integration gives access to a common behavioral measurement. Each downstream application can then be built as a lightweight, independent and specialized layer. These layers are not the measurement engine itself — they consume its signal.



One measurement signal. Multiple applications. Multiple infrastructures.

A signal is not a verdict.

It is an operational input for your organization's interpretation and decision.

- 1 Observation during execution
- 2 Production of measurements
- 3 Signal detection
- 4 Interpretation by your organization's policy
- 5 Decision and action governed by your organization
- 6 Creation of a verifiable trace

continue

verify

regenerate

reroute

escalate

validate

limit

stop

trace

NeoMundi measures and signals. Your organization interprets, governs and acts.

Observe what can still be corrected. Govern what cannot.

OBS — Governance observability

For responses or outputs that can still be corrected.

- Baseline
- Longitudinal tracking
- Comparison
- Alerts
- History
- Trace
- Drift analysis
- Semantic analysis

Full privacy. NeoMundi can receive only the necessary metrics, signals, snapshots or artefacts. Prompts and responses can stay within your perimeter depending on the architecture chosen.

GOV — Governance during execution

For actions that are hard to correct after execution.

- Continue
- Regenerate
- Reroute
- Validate
- Limit
- Stop
- Escalate

No retention. Prompts can pass through in real time with no durable storage logic. Retention depends on the integration mode and your policy.

Central criterion: can the action still be corrected after execution?

Your infrastructures stay in place.

NeoMundi does not replace your monitoring, compliance, audit or orchestration tools. It provides them with a common behavioral measurement that improves the quality of their observations, interpretations and decisions.

Can connect with it

GRC platforms, compliance platforms, observability tools, auditors, certifiers, insurers, risk teams, AI integrators, agentic systems, orchestration platforms.

Technical consumption

One API call, with no reconfiguration of existing infrastructure. BYOK-compatible where applicable. Data minimization; no storage of prompts and responses in the relevant modes.

NeoMundi does not replace your governance: it gives it a common signal, usable by the tools you already have.

Who for, and for what

Financial and fiduciary decisions

Banking, insurance, credit, fraud, compliance, transaction validation.

Autonomous and multi-agent systems

Agents connected to tools, APIs, databases and third-party systems.

Risk, compliance and audit teams

Access to actionable technical signals and artefacts for their processes.

Sensitive professional workflows

Healthcare, legal, public sector, human resources, critical services.

AI integrators and platforms

A governance, oversight and proof layer they can embed in their own offerings.

Regulatory compliance

Signals and traces that can support the EU AI Act (Art. 9, 12, 14, 26, 72, 73) and GDPR — without replacing legal analysis, certification or the competent authority.

Start where an error has a measurable consequence.

An instrument backed by continuous observation

The NeoMundi Observatory documents AI behaviors over time and feeds ControlTower's maturity. It publishes weekly barometers, behavioral cartographies, protocols, named studies, longitudinal analyses, methodological reports and de-identified public data. Its purpose is not to produce a commercial ranking — it is to make visible behaviors that would otherwise remain unmeasured.

FLAG signals correlated with confirmed problematic outputs in around 76 % of cases across a cumulative corpus.
Result observed under the conditions studied, not a universal guarantee.

Each observation improves the instrument, enriches the knowledge base and strengthens the ability to govern AI over time.

Infomaniak — The NeoMundi Observatory is developed with the support of Infomaniak, a Swiss infrastructure partner committed to digital sovereignty and a responsible cloud.

NVIDIA Inception — NeoMundi is a member of the NVIDIA Inception program, dedicated to startups developing advanced artificial intelligence technologies.

One workflow. One baseline. One measured value.

The pilot measures value instead of promising it abstractly.

- 1 Select a workflow with real stakes.
- 2 Establish the behavioral baseline.
- 3 Identify useful signals.
- 4 Define the intervention rules.
- 5 Measure operational value.

Human oversight

Human time per case, escalation rate, share of genuinely useful reviews.

Runtime costs

Tokens generated, avoidable regenerations, latency, cost per acceptable output.

Detection

Time between signal and intervention, earlier drift detection.

Audit and proof

Reconstitution time, artefact quality, proof availability.

Pilot output: baseline, observed signals, intervention rules and measured impact.

What is available today

A clear distinction between what is operational, what is being built in pilots with early clients, and what continues to be industrialized.

AVAILABLE TODAY

The measurement layer, usable right now

Runtime measurement layer and interfaces, API access, structured and timestamped signals. Sandbox, integration and interoperability documentation, measurement contract. Existing corpus of observations and experimental work.

PILOTS AND OPERATIONAL ARTICULATIONS

What is being built with early clients

Agents, orchestrators, governance systems, evidence infrastructures, business applications and partner integrations.

INDUSTRIALIZATION IN PROGRESS

What continues to be consolidated

Extending the measurement domain, metrological consolidation, further metric validation, performance, interoperability standardization, deployment modes, deepening potential predictive signals.

WHY NOW

AI no longer just answers. It acts.

Tool calls, workflow triggers, agent coordination, communications initiated with limited human oversight. The risk shifts from the final response to the authorization to act — and that authorization requires a behavioral context measured continuously, not just at deployment time.

Measure AI behavior.

Feed every use case.

Connect every infrastructure.

NeoMundi — The missing measurement context for AI behavior at runtime.

contact@neomundi.io · controltower.neomundi.io · neomundi.io