

NeoMundi — Executive Brief 2026

Continuous control of AI response risk

Reduce risk. Improve operational control. Produce the audit trail.

NeoMundi has developed **ControlTower**, a diagnostic, continuous control and traceability instrument designed to manage the risk associated with responses generated by artificial intelligence systems.

Privacy-first, ControlTower integrates through a simple API call, without requiring any infrastructure changes. The platform can also be deployed in sovereign mode.

NeoMundi measures. Your systems decide. Your auditors have the trail.

A level of measurement adapted to criticality

OBS — Observe without exposing content

Your prompts and responses remain within your environment.

NeoMundi receives only the normalized metrics and technical artefacts required to observe risk.

Never the prompt. Never the response.

GOV — Govern when intervention becomes necessary

When intervention during generation becomes critical, NeoMundi orchestrates the flow.

ControlTower measures how the generation regime evolves in real time, token by token, enabling the application of governance rules adapted to the level of risk.

The content transits temporarily. It is not retained, logged or stored.

Which systems is it designed for?

ControlTower is particularly relevant for organizations operating:

- AI systems in production;

- autonomous agents;
- sensitive workflows;
- high-stakes business applications;
- use cases subject to enhanced traceability, auditability or human-oversight requirements.

One API call. Runtime measurement. Your governance logic remains under your control.

The problem

AI responds. Your organization remains accountable.

A response has been generated.

But under what behavioural conditions was it produced?

An unobserved drift.

An atypical regime that went undetected.

A fluent, plausible output that is insufficiently reliable.

No trail when an audit becomes necessary.

In many AI deployments, teams know that a generation occurred.

They do not always have an actionable signal to:

- observe runtime behaviour;
- detect certain drifts;
- identify responses requiring closer scrutiny;
- document generation conditions;
- trigger human review;
- demonstrate months later that controls were actually in place.

Every insufficiently documented generation becomes an operational risk that is difficult to assess, supervise and audit.

What NeoMundi does

NeoMundi adds a complementary layer of behavioural measurement and governance to generative and agentic AI systems.

ControlTower produces observable signals and timestamped artefacts that your systems can:

- classify;
- monitor;
- route;
- escalate;
- block according to the selected policy;
- integrate into their own governance workflows.

NeoMundi does not claim to determine the semantic truth of responses on its own.

NeoMundi does not decide on your behalf.

The instrument measures, signals and produces an actionable trail. Your governance logic remains under your control.

Observable evidence

An AI system can be stable and still be wrong.

Traditional benchmarks mainly evaluate the final response.

NeoMundi also observes the behavioural conditions under which that response is generated.

This distinction reveals certain regimes that would remain invisible in a simple output evaluation.

Initial experimental results

NeoMundi conducted two controlled campaigns on generative AI services accessible through APIs.

Combined scope

Result

Generations analysed **10,160**

FLAGS triggered **≈ 831**

Problematic outputs confirmed among FLAGS **≈ 632**

Observed precision of the FLAG signal **≈ 76%**

When a FLAG was triggered, a problematic output was confirmed in approximately **76% of cases** across the combined corpus.

The protocol compares the signals generated by NeoMundi with a complementary verification procedure applied to the observed outputs. The methodology is documented separately.

The FLAG captures a useful subset of problematic responses.

It does not cover every error: some errors occur under an apparently stable regime.

This is why NeoMundi does not reduce governance to a single signal. The objective is to combine several dimensions of measurement to better observe AI behaviour in production.

NeoMundi does not claim to detect every error.

The instrument currently prioritizes signal precision over exhaustive coverage: it is better to flag fewer responses, but flag them usefully, than to overwhelm teams with false positives.

These results should be interpreted within their limitations: they depend on the corpus, the providers tested, the selected thresholds and the confirmation protocol.

Consolidation continues through observation campaigns, methodological audits and field pilots.

Two deployment modes

Observe continuously. Govern when intervention becomes critical.

ControlTower adapts to the risk level of your system.

	OBS — Governance observability	GOV — Runtime orchestration
Purpose	Continuous diagnostics and observability	Control during execution
Timing	After generation or outside the critical path	During generation
Data processed	Normalized metrics and technical artefacts	Temporarily processed flow
Confidentiality	No prompt or response transmitted	Temporary transit, without content retention
Integration	Your pipeline transmits the expected artefacts	NeoMundi orchestrates the call to the LLM provider
Signals	Stability status, alerts and auditable trails	Stability status, ΔG , alerts and governance decisions
Natural use case	The majority of AI systems	Critical workflows or outputs that are difficult to correct

The right decision criterion

If an incorrect output reaches the user, can it still be corrected?

Yes

An alert, downstream moderation, regeneration, a correction during the next cycle or human review remain possible.

OBS is generally the natural starting point.

No

The response may concern, for example, medical advice, a legal decision communicated directly, a fiduciary decision, critical infrastructure or an autonomous agent with the ability to act.

GOV becomes relevant.

In both cases, NeoMundi provides the signal.

Your system, your rules layer or the responsible operator retains decision-making authority.

Traceability, auditability and evidence-based governance

ControlTower associates observed responses with observation identifiers, calculated metrics, governance scores, risk signals, timestamps and configurable traceability artefacts.

The objective is not to ask your teams to trust a promise.

The objective is to provide them with an actionable trail.

Contribution to AI Act and GDPR compliance programmes

ControlTower does not replace a comprehensive compliance programme, legal analysis or regulatory certification.

The platform nevertheless provides several directly useful technical capabilities:

- continuous monitoring of AI system behaviour;
- control of the risk associated with generated responses;
- operational traceability and auditability;
- actionable signals for human oversight;
- evidence useful for incident documentation;
- a privacy-first architecture based on data minimization.

ControlTower is designed around three principles:

zero content retention · zero content logging · zero content storage

Minimal measurement and traceability artefacts may remain, depending on the selected mode and configuration.

Start simply

For most organizations, OBS is the natural starting point.

One API call. No infrastructure changes. No prompt or response transmitted.

Start with a targeted flow, observe the signals and progressively calibrate your governance policy.

When the level of criticality requires it, move to GOV.

Test ControlTower

[Access the sandbox](#) · [Discover NeoMundi](#) · [Consult the AI Act and GDPR mapping](#)

Measure what can be corrected. Control what cannot.

contact@neomundi.io